

Memory Allocation and Scope

Lecture 4

Christian A. Pagot



Universidade Federal da Paraíba
Centro de Informática

Physical Memory (RAM)

- May not be **sufficient** to fulfill our needs.
- Certain **code** portions may be **rarely used**.
- Physical memory should be **abstracted**.
- Usually, processes **should not interfere** with each other.

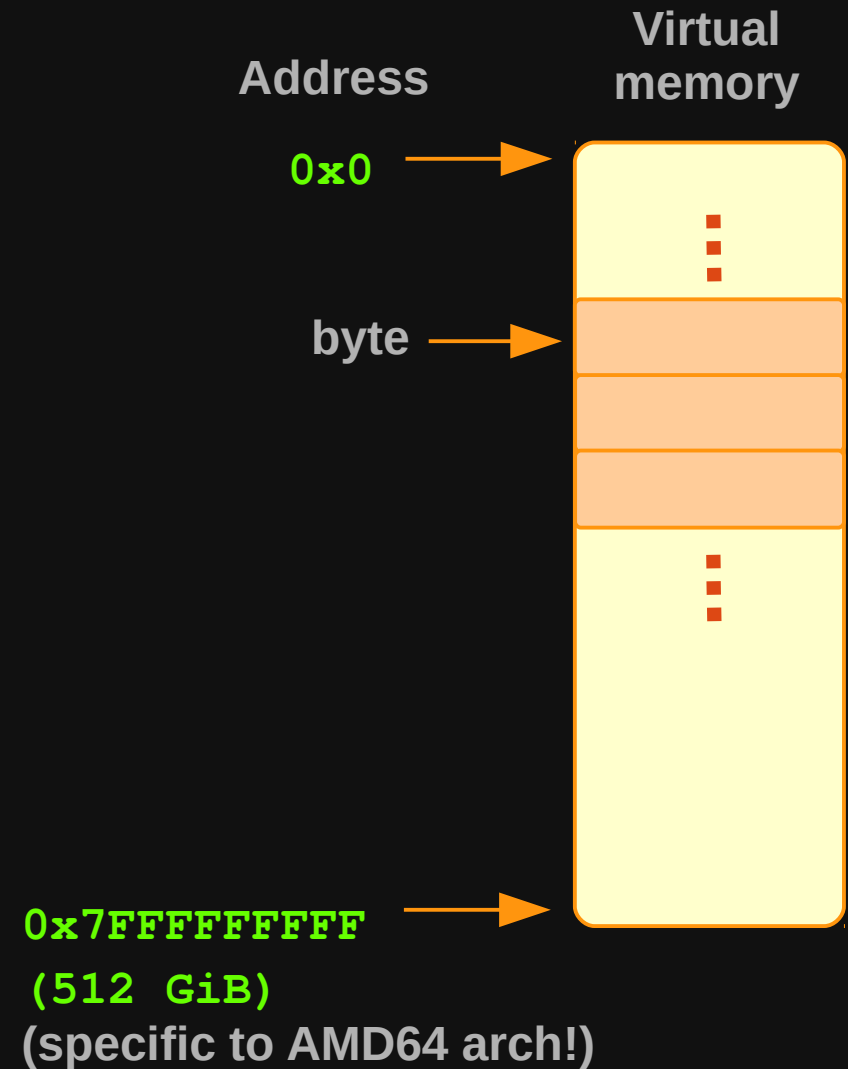
How we could
handle these
problems/limitations?

Virtual Memory!



Virtual Memory

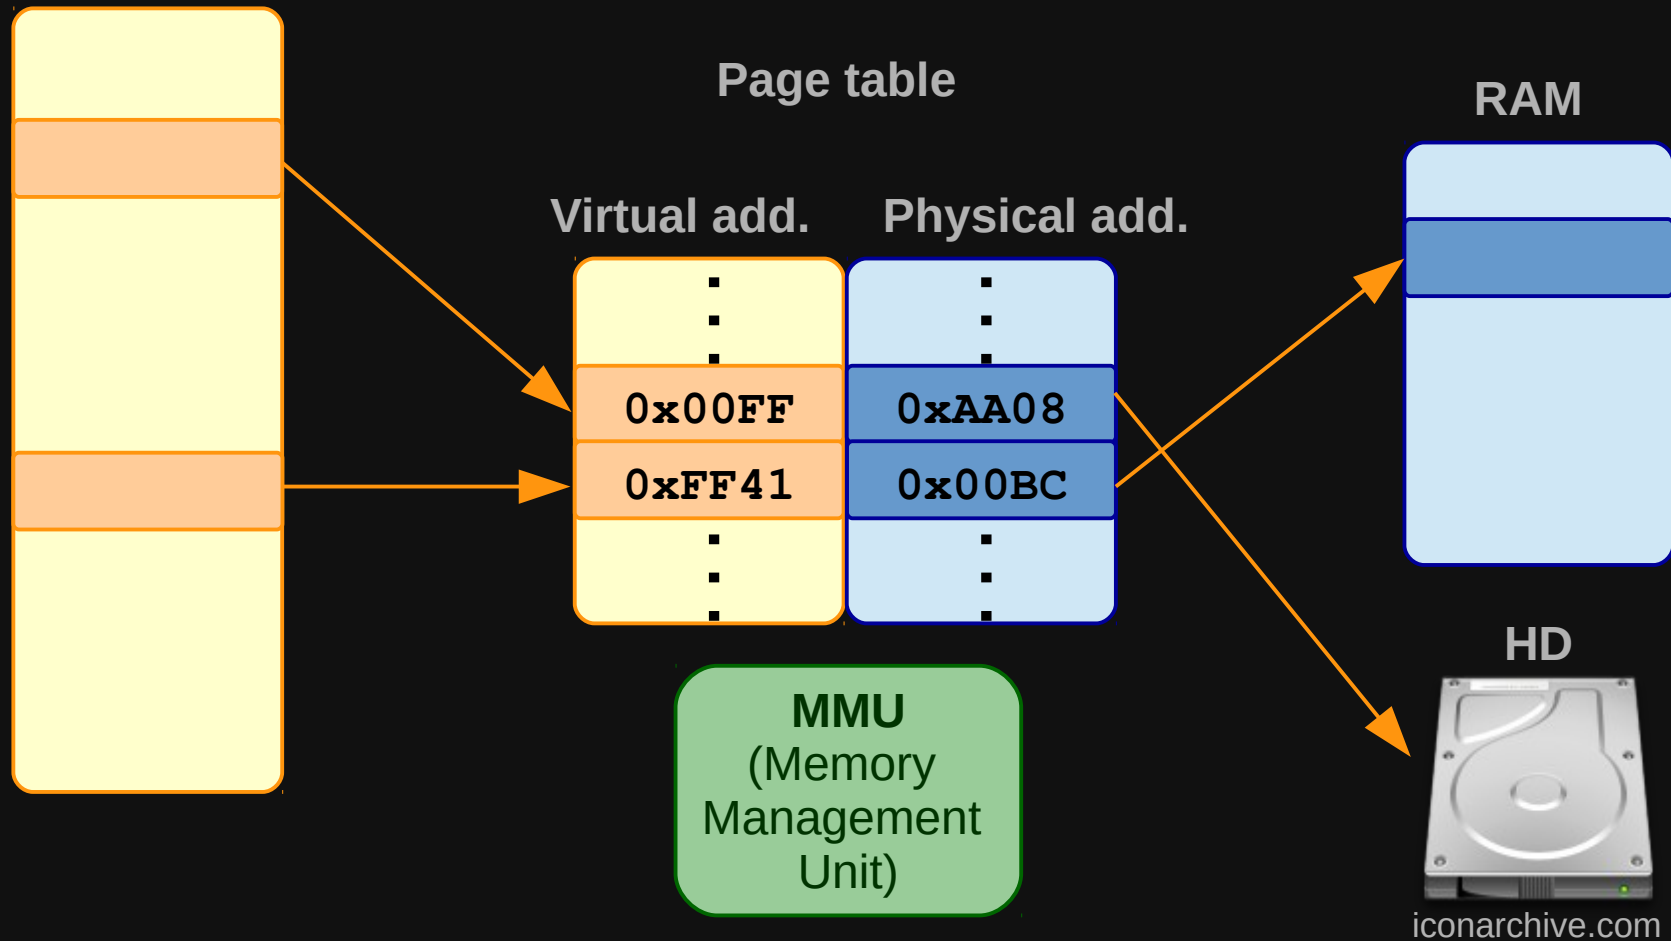
- Main storage appears as a **continuous** address space.
- Maps **virtual** addresses to **physical** addresses.
- The OS **manages** virtual address spaces.
- Address translation is accomplished by the CPU, through dedicated hardware: **Memory Management Unit (MMU)**.
- Even **files** (located on the HD) can be **mapped** into the virtual memory!



Virtual Memory

Virtual memory

Physical memory




Virtual Memory

Program A
(virtual mem.)



Program B
(virtual mem.)



Page table for prog. A

Virtual add.	Physical add.
⋮	⋮
0x00AA	0xCF11
⋮	⋮

Page table for prog. B

Virtual add.	Physical add.
⋮	⋮
0x00AA	0xEE20
⋮	⋮

Physical memory



How a **program** is
organized within the
virtual memory?



Physical / Virtual Memory Example

Checking current physical / virtual memory usage

```
~$ man top
```

```
~$ top
top - 16:14:06 up 2:19, 2 users, load average: 0,57, 0,52, 0,46
Tasks: 232 total, 1 running, 231 sleeping, 0 stopped, 0 zombie
%Cpu(s): 7,8 us, 2,0 sy, 0,0 ni, 90,0 id, 0,0 wa, 0,0 hi, 0,2 si, 0,0 st
KiB Mem: 3949960 total, 2980252 used, 969708 free, 96540 buffers
KiB Swap: 22142972 total, 0 used, 22142972 free. 1168940 cached Mem

  PID USER      PR  NI    VIRT    RES    SHR S  %CPU  %MEM     TIME+ COMMAND
 1323 root        20   0   250052   84100   46968 S   9,6   2,1   8:10.86 Xorg
 2371 christi+   20   0  1267512  246364  133924 S   9,6   6,2   7:18.72 chrome
 3272 christi+   20   0  1044068  232076   77696 S   7,6   5,9   4:11.55 chrome
 2082 christi+   20   0  1390992  191244   84784 S   6,0   4,8   3:18.66 compiz
```

VIRT : virtual memory size (code, data, shared libraries, etc.).

RES : resident memory size (non-swapped physical memory).

SWAP : swapped size (non-resident portion of the task's address space).

SHR : memory potentially shared with other processes.

top: press 'f' to
select columns!

DIY!



Storage Class and Scope

- Programs are composed of:
 - Executable code.
 - Data.
- Data
 - Pertains to a storage class, which defines its lifespan.
 - Has a scope, which defines its visibility within the program.
 - *“Storage class and scope are assumed from the location of a datum's declaration, and determine its placement within virtual memory.”*



Storage Class and Scope

- Data declared **outside** any **function** has
 - Global scope.
 - Static duration (it exists for the program lifespan).
- Example

example_31.c

```
int x = 7;

void f() {
    x = 19;
}

int main() {

    x = 13;
    f();

    return 0;
}
```

- **x** has **global scope**.
- **x** has **static duration**.



Storage Class and Scope

- Data declared **inside** a function
 - Local scope.
 - Automatic duration (it exists for the duration of the function call).
- Example

`example_32.c`

```
void f() {  
    int x = 19;  
    x++;  
}  
  
int main() {  
  
    f();  
    f();  
  
    return 0;  
}
```

- **x** has **local scope**.
- **x** has **automatic duration**.



Programs vs. Processes

- Program

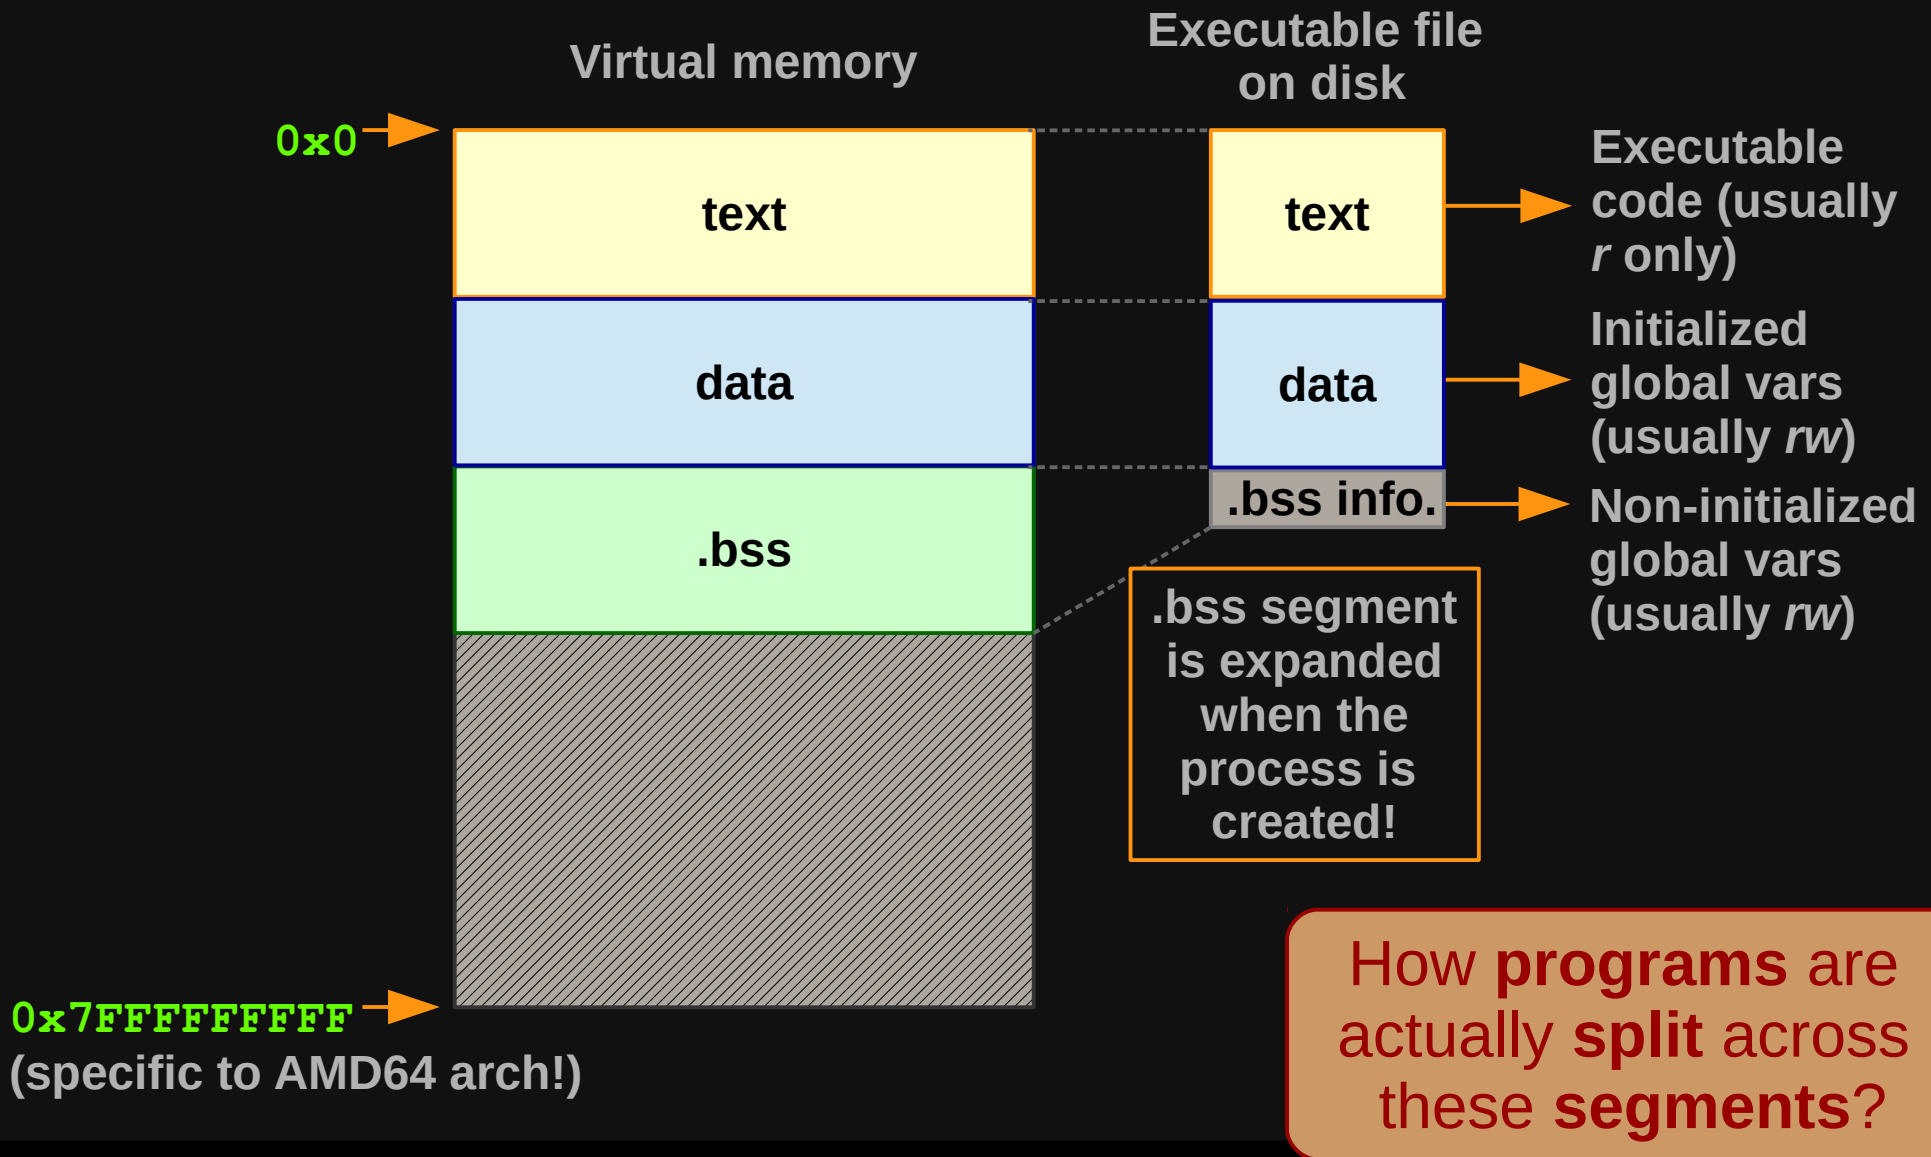
- Source file.
- Compiled source file on disk.

- Process

- “*Instance of a computer program that is being executed*” (Wikipedia).
- Consists of:
 - PID (Process IDentification number).
 - Context (register values).
 - Page table.
 - List of open files, etc.

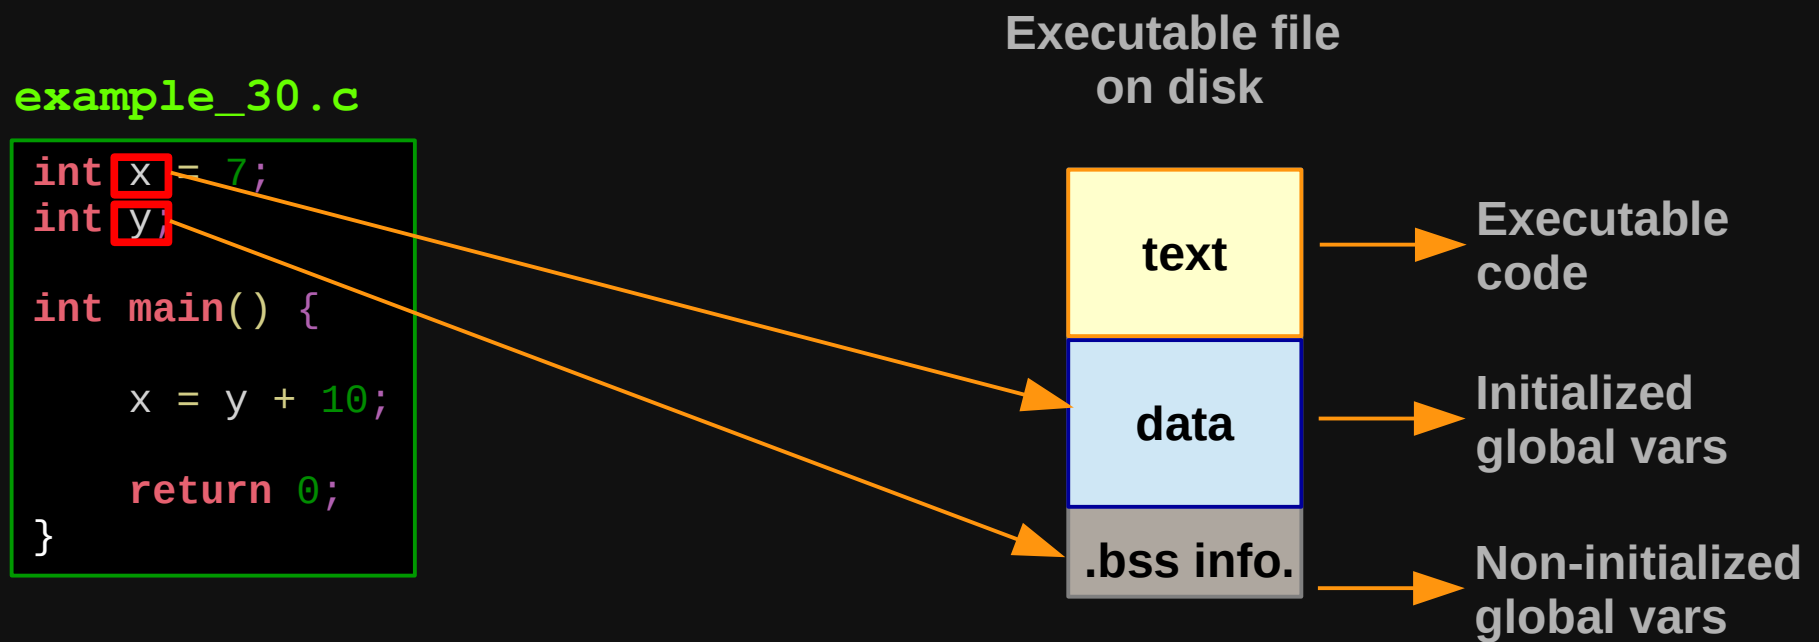


Process Memory Layout



Program Segmentation Example

- **Executable and Linking Format (ELF)** is the executable file format used by **Linux**.
- **Data and instructions** are aggregated by the compiler in distinct **segments**.



Process Memory Layout Example

Example

example_34.c

```
char x[104857600];  
  
int main (void)  
{  
    return 0;  
}
```

DIY!

```
~$ gcc example_34.c -o example_34
```

```
~$ ls -l
```

```
-rwxrwxr-x 1 christian christian 8532 Jan 29 20:30 example_34  
-rw-rw-r-- 1 christian christian 55 Jan 29 20:29 example_34.c
```

```
~$ man size
```

```
~$ size example_34
```

text	data	bss	dec	hex	filename
1115	552	104857632	104859299	64006a3	example_34



More on Storage Class and Scope!

- Data with **local scope** can present **static duration** if declared with the **static** qualifier.
 - It is important when we want to retain the value of a variable between function calls!
- Example

example_33.c

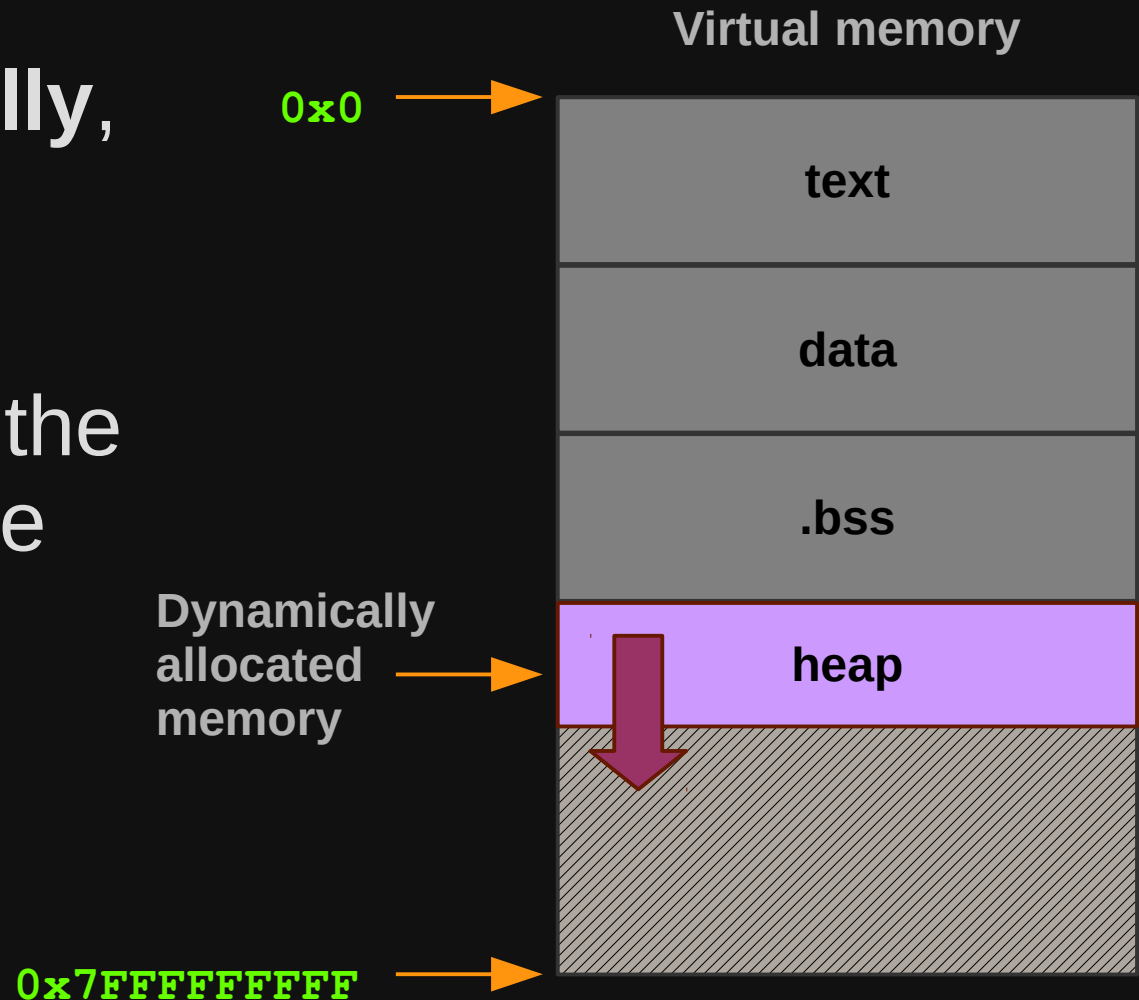
```
void f() {  
    static int x = 19;  
    x++;  
}  
  
int main() {  
  
    f();  
    f();  
  
    return 0;  
}
```

- **x** has **local scope**.
- **x** has **static duration**.



Dynamic Memory Allocation

- Memory can be allocated **dynamically**, during **runtime**.
- Dynamic memory is allocated right **after** the **data segment**, in the **heap segment**.



Dynamic Memory Allocation

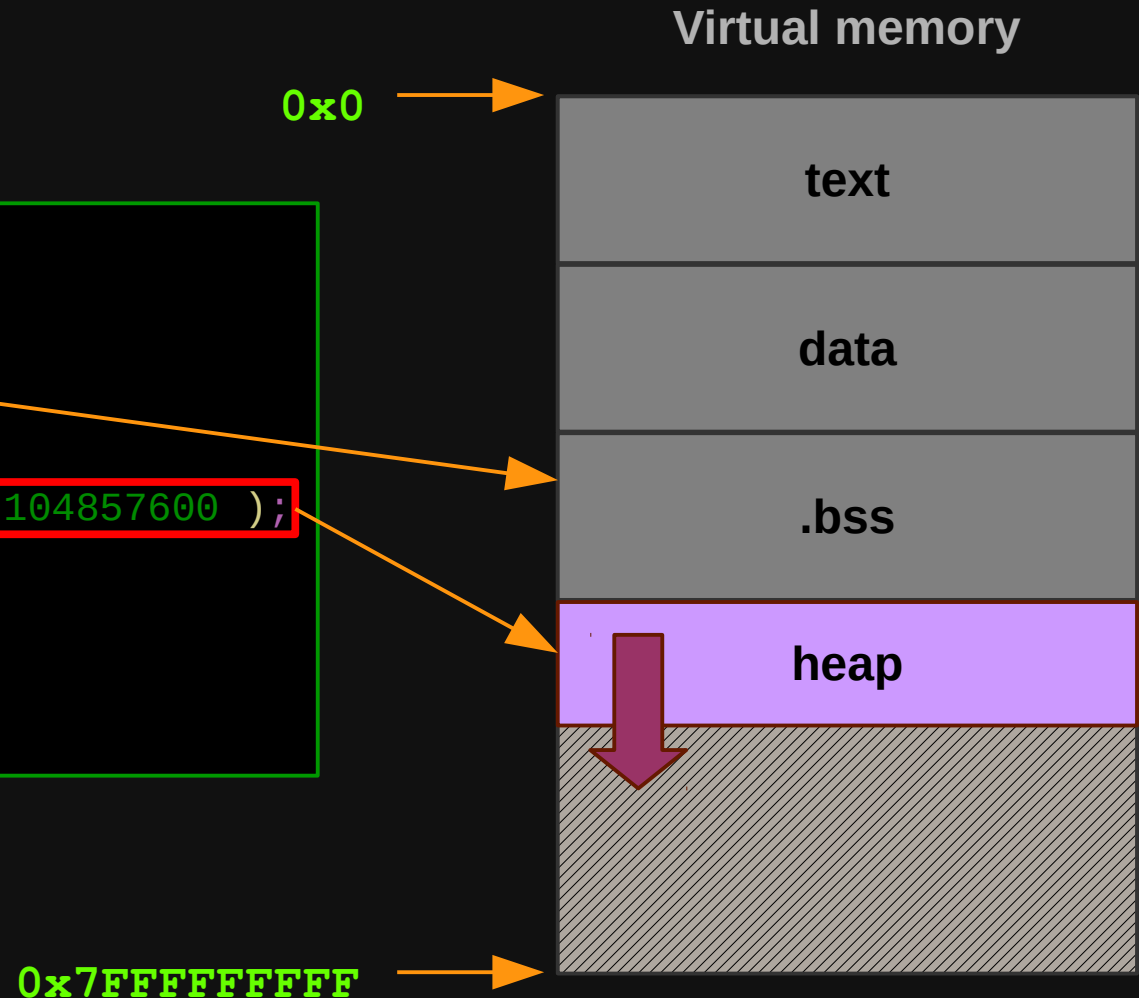
Example

example_35.c

```
#include<stdlib.h>

char * pc;

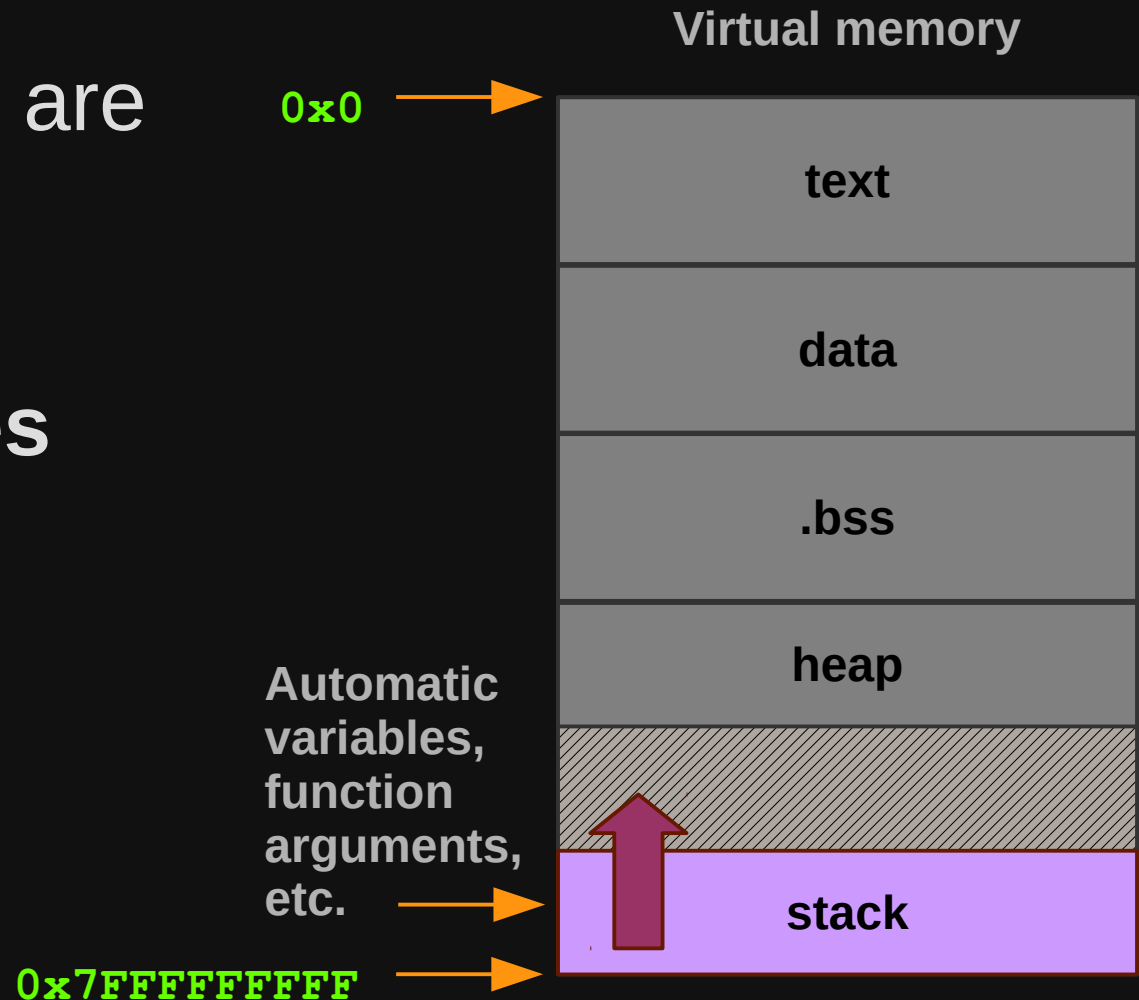
int main (void)
{
    pc = malloc( sizeof( char ) * 104857600 );
    free( pc );
    return 0;
}
```



Automatic Variables

- **Local variables and function arguments are allocated automatically.**
- **Automatic variables are pushed onto the stack.**

More on **parameter passing** later!



Automatic Variables

Example

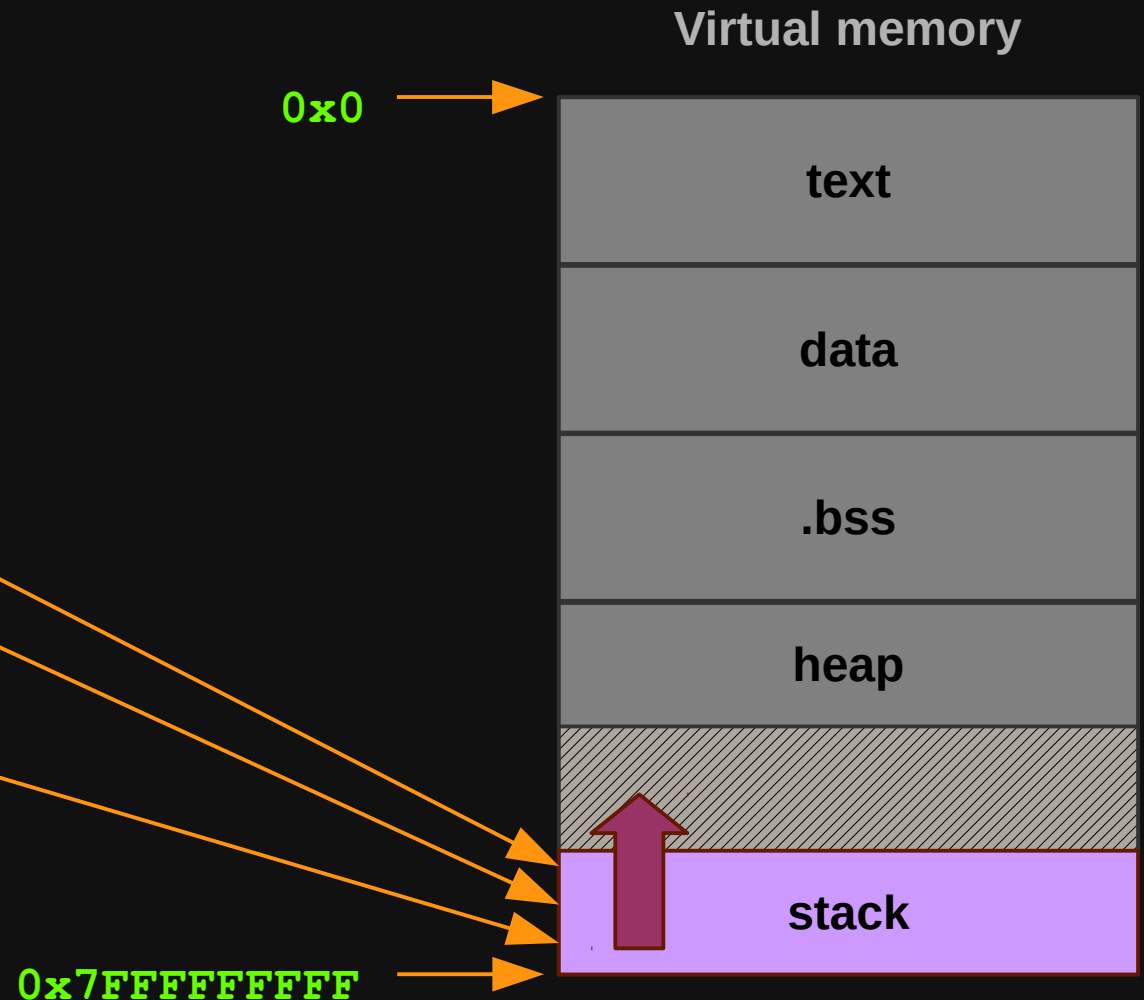
example_37.c

```
#include <stdio.h>

void f( int i ) {
    int y = i + 20;
    printf( "value: %i\n", y );
}

int main() {
    int x = 10;
    f( x );

    return 0;
}
```



Stack Facts

- GNU compilers allow memory to be allocated **dynamically** on the current **stack frame** with **`alloca()`**.

There is no need to explicitly free **`alloca()`**'ed memory. It is automatically released at the end of the current function call.

- Stacks usually have a **maximum** prescribed **size**.

Thus, take care when using **`alloca()`**!

- Stacks allow for **recursion**!



More on “*the terrible story of `alloca()` and the `inline` function*” later!



C99 Variable Length Array (VLA)

- It's an **array type** whose **size** is determined only at **runtime**.

Exercise:

- Try using GDB to find how the VLA is allocated on your system.
- Could you describe at least one difference between the VLA and **alloca()**?



Parameter Passing

- C supports **by-value** parameter passing!

Everything passed to functions are copies, including pointers!

- Example

byvalue.c

```
#include <stdio.h>

void f( int i ) {
    i = i + 1;
    printf( "value: %i\n", i );
}

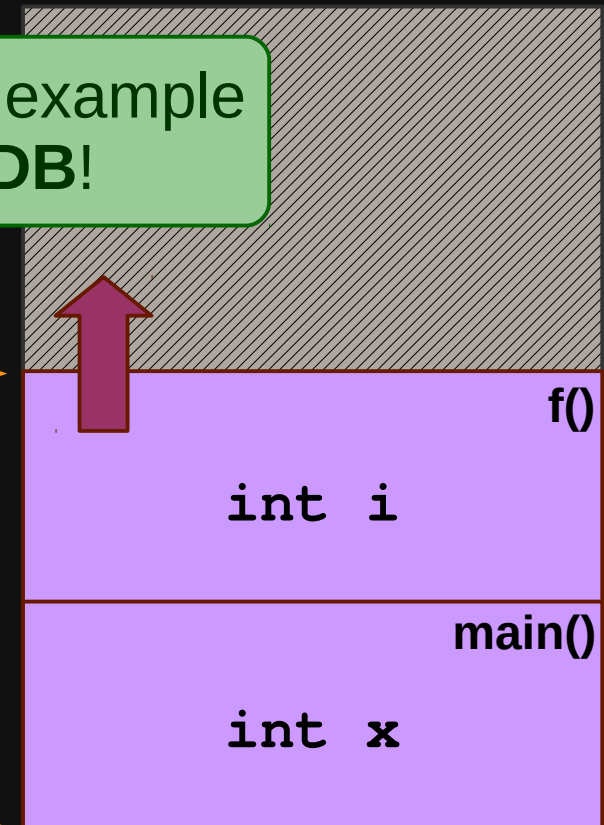
int main() {
    int x = 10;
    f( x );
    printf( "value: %i\n", x );

    return 0;
}
```

Explore this example
with **GDB**!

Top of the
stack

0x7FFFFFFF



Parameter Passing

- And how about **by-reference** parameter passing?

Actually, it is by-value!!

- Example

byreference.c

```
#include <stdio.h>

void f( int *i ) {
    *i= *i + 1;
    printf( "value: %i\n", *i );
}

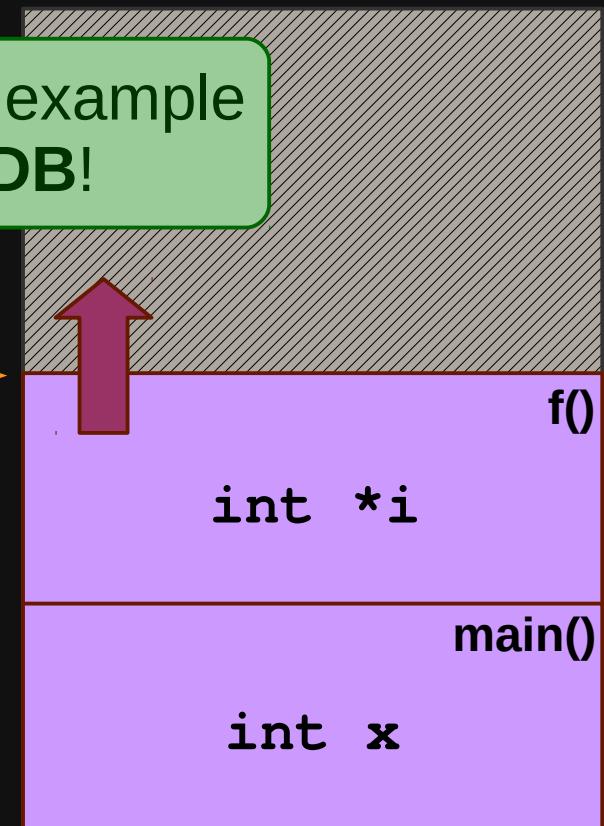
int main() {
    int x = 10;
    f( &x );
    printf( "value: %i\n", x );

    return 0;
}
```

Explore this example
with **GDB**!

Top of the
stack

0x7FFFFFFF



Changing a Pointer Parameter

Inserting a new node after the given one:

ptr2struct.c

```
#include <stdio.h>
#include <stdlib.h>

struct Node {
    int value;
    struct Node *next;
};

void InsertAfter(struct Node *root, int value) {
    root->next = (struct Node*) malloc (sizeof(struct Node));
    root->next->value = value;
    root->next->next = NULL;
}

int main() {
    struct Node *n1 = (struct Node*) malloc (sizeof(struct Node));
    n1->value = 10;
    n1->next = NULL;
    InsertAfter(n1, 20);

    return 0;
}
```

Implement the function
InsertBefore(...)!

→ Given a pointer to the first node, insert the new node before it. Update the pointer parameter to point to the new node!

DIY!



Changing a Pointer Parameter

Solution:

ptr2ptr2struct.c

```
#include <stdio.h>
#include <stdlib.h>

struct Node {
    int value;
    struct Node *next;
};

void InsertBefore(struct Node **root, int value) {
    struct Node *tmp_ptr = (struct Node*) malloc (sizeof(struct Node));
    tmp_ptr->value = value;
    tmp_ptr->next = *root;
    *root = tmp_ptr;
}

int main() {
    struct Node *n1 = (struct Node*) malloc (sizeof(struct Node));
    n1->value = 10;
    n1->next = NULL;
    InsertBefore(n1, 20);

    return 0;
}
```

You will need a **pointer to a pointer** in order to **change the pointers' value!**



References

- **Understanding Memory.** University of Alberta, 2008.
- **Virtual Memory.** Wikipedia.
 - https://en.wikipedia.org/wiki/Virtual_memory
- **Memory.** Florent Bruneau. Intersec TechTalk. 2013.
 - <https://techtalk.intersec.com/2013/07/memory-part-1-memory-types>
- **Linux Memory Management System Source Code**
 - <http://git.kernel.org/cgit/linux/kernel/git/torvalds/linux.git/tree/mm>



Further Reading

- **What Every Programmer Should Know About Memory.** Ulrich Drepper. Red Hat, Inc. 2007.

