

Looking at the Assembly

Lecture 5

Christian A. Pagot



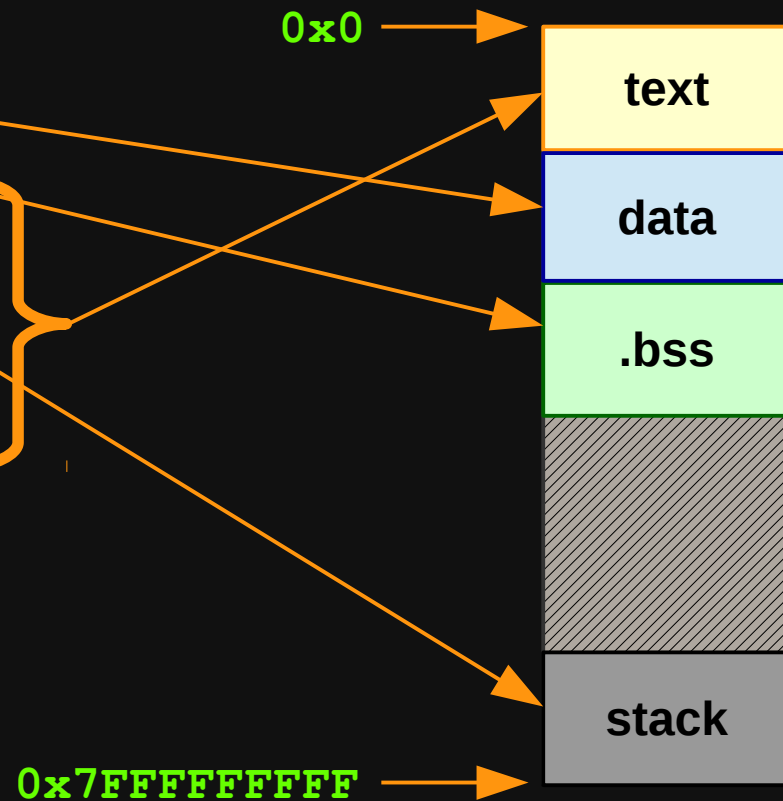
Universidade Federal da Paraíba
Centro de Informática

Data and Code Segmentation

How the **elements** of the program bellow are actually **segmented** within the **virtual memory**?

assembly1.c

```
int x = 10;  
int y  
  
main() {  
    int w  
    y = 20;  
    w = x + y;  
    return 0;  
}
```



Check the **addresses** of the **variables** and the **main function**!

For a less abstract view, let's see what is happening at the **assembly** level...



Assembly

*“An **assembly language** (or assembler language) is a **low-level** programming language for a computer, or other programmable device, in which there is a very strong (generally one-to-one) **correspondence** between the **language** and the architecture's **machine code instructions**.”*

Assembly language, Wikipedia



x86 and x86-64 (AMD64)


- x86
 - Intel's and AMD's 32 bit instruction set.
- x86-64 or AMD64
 - Intel's and AMD's 64 bit instruction set.

We will keep with the
x86-64 instruction set!

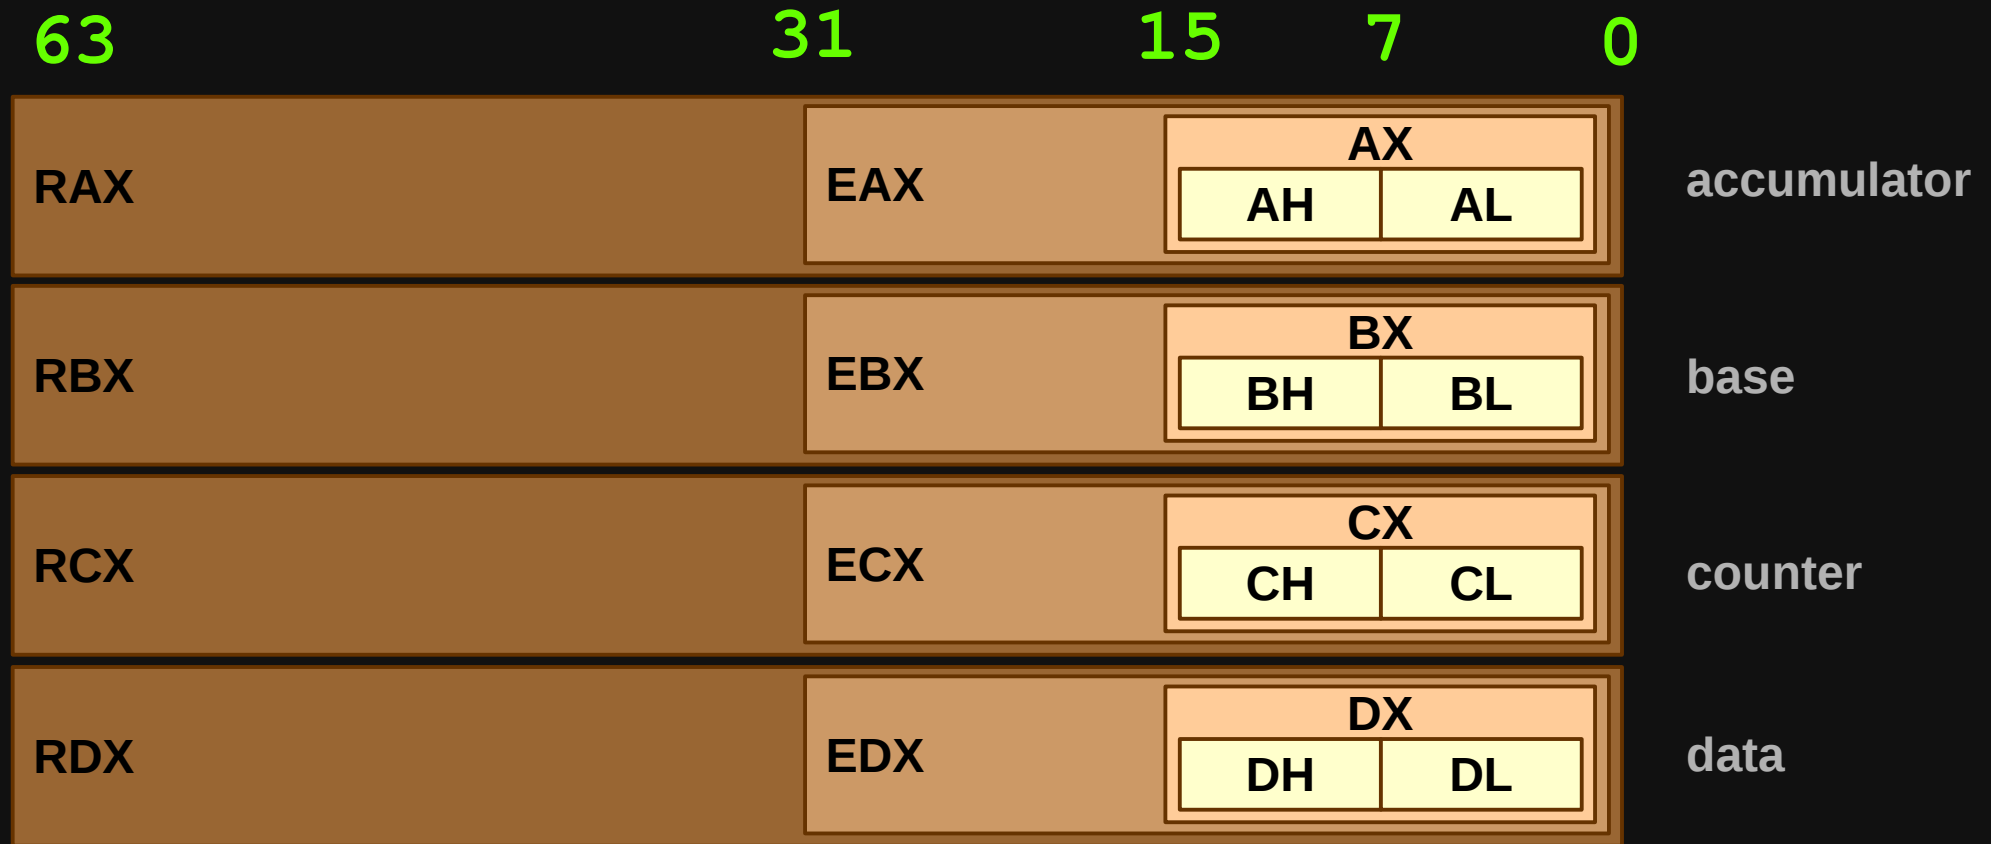


x86-64 Register Set

Register widths:



16 General Purpose Register Set



16 General Purpose Register Set

63	31	15	7	0	
RBP	EBP	BP	BPL		stack base pointer
RSI	ESI	SI	SIL		source index
RDI	EDI	DI	DIL		destination index
RSP	ESP	SP	SPL		stack pointer



16 General Purpose Register Set

63	31	15	7	0
R8	R8D	R8W	R8B	
R9	R9D	R9W	R9B	
R10	R10D	R10W	R10B	
R11	R11D	R11W	R11B	
R12	R12D	R12W	R12B	
R13	R13D	R13W	R13B	
R14	R14D	R14W	R14B	
R15	R15D	R15W	R15B	



Other Registers

- **RIP** : 64 bit instruction pointer.
 - Points to the next instruction to be executed.
- **RFLAGS** : Flags.
 - Stores flags generated by computation results and for controlling the processor.
- **FPR0, ..., FPR7** : Floating Point Unit (FPU) registers.
 - Floating point registers.
- **MMX and XMM**
 - A set of registers to be used by the MMX and SSE extensions.



Assembly Syntax

- There are two assembly syntaxes that are very **popular**:
 - Intel.
 - AT&T.

We will keep with the
AT&T syntax!



AT&T Assembly Syntax

- General format of instructions
 - Mnemonic source, destination
- Operation suffixes:

b	byte (8 bit)
s	short (16 bit integer) or single (32-bit floating point)
w	word (16 bit)
l	long (32 bit integer or 64 bit floating point)
q	quad (64 bit)
t	ten bytes (80-bit floating point)

- Prefixes
 - Registers are prefixed with %.
- Example
 - **movb** \$0x05, %al



AT&T Assembly Syntax

- Address operand syntax

`segment:displacement(base register, offset register, scalar multiplier)`

- Examples

- `movl -4(%ebp, %edx, 4), %eax`
 - Full example: load $*(\text{ebp} - 4 + (\text{edx} * 4))$ into `eax`.
- `movl -4(%ebp), %eax`
 - Typical example: load a stack variable into `eax`.
- `movl (%ecx), %edx`
 - No offset: copy the target of a pointer into a register.



Obtaining the Assembly of C Code

- **How to obtain the assembly code generated from C code?**
 - Typically, two methods are used:
 - Getting the output assembly code generated by the compiler.
 - Disassembling the machine code with a debugger during the runtime.



Assembly Code Generated by GCC

```
~$ gcc -S -fno-asynchronous-unwind-tables assembly1.c
```

assembly1.c

```
int x = 10;
int y;

main() {
    int w;
    y = 20;
    w = x + y;
    return 0;
}
```



assembly1.s

```
.file "disassembly1.c"
.globl x
.data
.align 4
.type x, @object
.size x, 4
x:
.long 10
.comm y, 4, 4
.text
.globl main
.type main, @function
main:
pushq %rbp
movq %rsp, %rbp
movl $20, y(%rip)
movl x(%rip), %edx
movl y(%rip), %eax
addl %edx, %eax
movl %eax, -4(%rbp)
movl $0, %eax
popq %rbp
ret
.size main, .-main
.ident "GCC: (Ubuntu 4.8.4-2ubuntu1~14.04) 4.8.4"
.section .note.GNU-stack,"",@progbits
```



GDB Disassembler

DIY!

```
(gdb) disassemble
```

assembly1.s

```
.text
.globl  main
.type   main,@function

main:
pushq   %rbp
movq    %rsp, %rbp
movl    $20, y(%rip)
movl    x(%rip), %edx
movl    y(%rip), %eax
addl    %edx, %eax
movl    %eax, -4(%rbp)
movl    $0, %eax
popq    %rbp
ret
```

Only the .text segment.

GDB output.

Dump of assembler code for function main:

```
0X00000000004004ed <+0>: push    %rbp
0X00000000004004ee <+1>: mov     %rsp,%rbp
=> 0x00000000004004f1 <+4>: movl    $0x14,0x200b45(%rip) # 0x601040 <y>
0X00000000004004fb <+14>: mov     0x200b37(%rip),%edx # 0x601038 <x>
0X0000000000400501 <+20>: mov     0x200b39(%rip),%eax # 0x601040 <y>
0X0000000000400507 <+26>: add     %edx,%eax
0X0000000000400509 <+28>: mov     %eax,-0x4(%rbp)
0X000000000040050c <+31>: mov     $0x0,%eax
0X0000000000400511 <+36>: pop     %rbp
0X0000000000400512 <+37>: retq
```

End of assembler dump.

Only the .text segment (actually, the main() function).

Symbols have been substituted by actual virtual memory addresses!



Disassembling with /m



Dump of assembler code for function main:

```
4  main() {
    0X0000000000004004ed <+0>: push    %rbp
    0X0000000000004004ee <+1>: mov     %rsp,%rbp

5      int w;
6      y = 20;
=> 0x0000000000004004f1 <+4>: movl    $0x14,0x200b45(%rip) # 0x601040 <y>

7      w = x + y;
    0X0000000000004004fb <+14>: mov     0x200b37(%rip),%edx # 0x601038 <x>
    0X000000000000400501 <+20>: mov     0x200b39(%rip),%eax # 0x601040 <y>
    0X000000000000400507 <+26>: add     %edx,%eax
    0X000000000000400509 <+28>: mov     %eax,-0x4(%rbp)

8      return 0;
    0X00000000000040050c <+31>: mov     $0x0,%eax

9  }
```

```
0X000000000000400511 <+36>: pop     %rbp
0X000000000000400512 <+37>: retq
```

End of assembler dump.



Inspecting the Assembly Code

```
4  main() {  
    0X000000000004004ed <+0>: push    %rbp  
    0X000000000004004ee <+1>: mov     %rsp, %rbp
```

Prologue or preamble.



Inspecting the Assembly Code

```
4  main() {  
   0X00000000004004ed <+0>: push    %rbp  
   0X00000000004004ee <+1>: mov     %rsp, %rbp
```

to the address
generated by this expression.
(which is the address of `y` in the
data segment!)

```
5      int w;  
6      y = 20;  
=> 0x00000000004004f1 <+4>: movl    $0x14, 0x200b45(%rip) # 0x601040 <y>
```

the value
move... 20...

```
7      w = x + y;  
   0X00000000004004fb <+14>: mov     0x200b37(%rip), %edx # 0x601038 <x>  
   0X0000000000400501 <+20>: mov     0x200b39(%rip), %eax # 0x601040 <y>  
   0X0000000000400507 <+26>: add     %edx, %eax  
   0X0000000000400509 <+28>: mov     %eax, -0x4(%rbp)
```

Store the resulting value at the computed
address (which is the address of `w`, on
the stack!)

Before we sum two
values, we have
to move them to registers...



Inspecting the Assembly Code

```
4  main() {  
    0X000000000004004ed <+0>: push    %rbp  
    0X000000000004004ee <+1>: mov     %rsp,%rbp
```

```
5      int w;  
6      y = 20;  
=> 0x000000000004004f1 <+4>: movl    $0x14,0x200b45(%rip) # 0x601040 <y>  
  
7      w = x + y;  
    0X000000000004004fb <+14>: mov     0x200b37(%rip),%edx # 0x601038 <x>  
    0X00000000000400501 <+20>: mov     0x200b39(%rip),%eax # 0x601040 <y>  
    0X00000000000400507 <+26>: add     %edx,%eax  
    0X00000000000400509 <+28>: mov     %eax,-0x4(%rbp)
```

```
8      return 0;  
    0X0000000000040050c <+31>: mov     $0x0,%eax
```

move...

the value
0...

to the EAX register
(which is usually used by the
functions to return values!)



Inspecting the Assembly Code

```
4  main() {  
   0X000000000004004ed <+0>: push    %rbp  
   0X000000000004004ee <+1>: mov     %rsp,%rbp
```

```
5      int w;  
6      y = 20;  
=> 0x000000000004004f1 <+4>: movl    $0x14,0x200b45(%rip) # 0x601040 <y>  
  
7      w = x + y;  
   0X000000000004004fb <+14>: mov     0x200b37(%rip),%edx # 0x601038 <x>  
   0X00000000000400501 <+20>: mov     0x200b39(%rip),%eax # 0x601040 <y>  
   0X00000000000400507 <+26>: add     %edx,%eax  
   0X00000000000400509 <+28>: mov     %eax,-0x4(%rbp)
```

```
8      return 0;  
   0X0000000000040050c <+31>: mov     $0x0,%eax
```

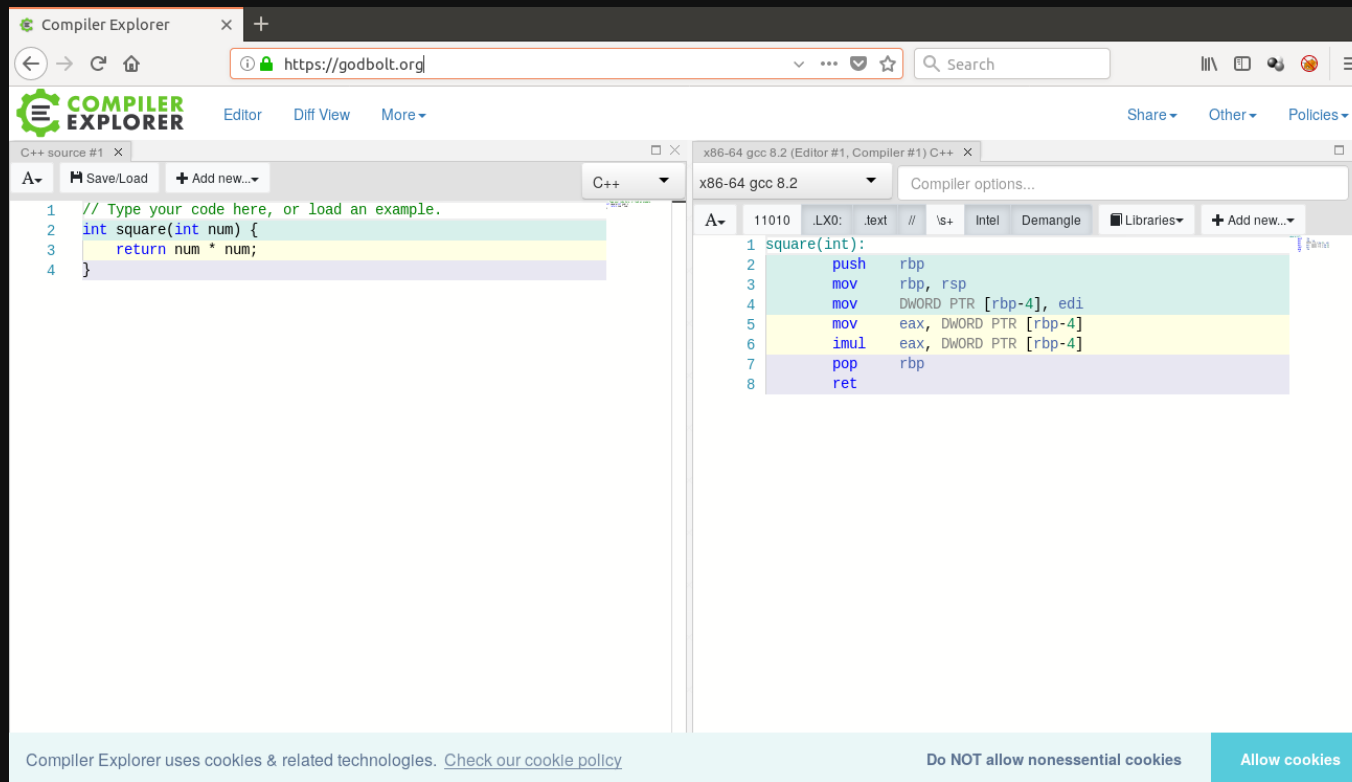
```
9  }  
   0X00000000000400511 <+36>: pop     %rbp  
   0X00000000000400512 <+37>: retq
```

Epilogue.



Inspecting the Assembly Code

- Cool C/D/C++/CUDA... to Assembly online service!
 - <https://godbolt.org>



The screenshot shows the Compiler Explorer interface. On the left, the C++ source code is displayed:

```
1 // Type your code here, or load an example.
2 int square(int num) {
3     return num * num;
4 }
```

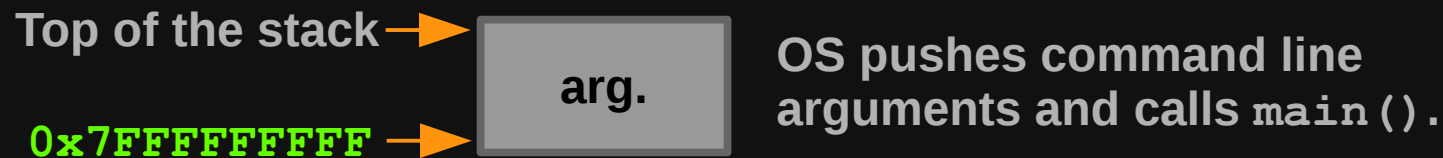
 On the right, the assembly code generated by x86-64 gcc 8.2 is shown:

```
1 square(int):
2     push    rbp
3     mov     rbp, rsp
4     mov     DWORD PTR [rbp-4], edi
5     mov     eax, DWORD PTR [rbp-4]
6     imul    eax, DWORD PTR [rbp-4]
7     pop     rbp
8     ret
```

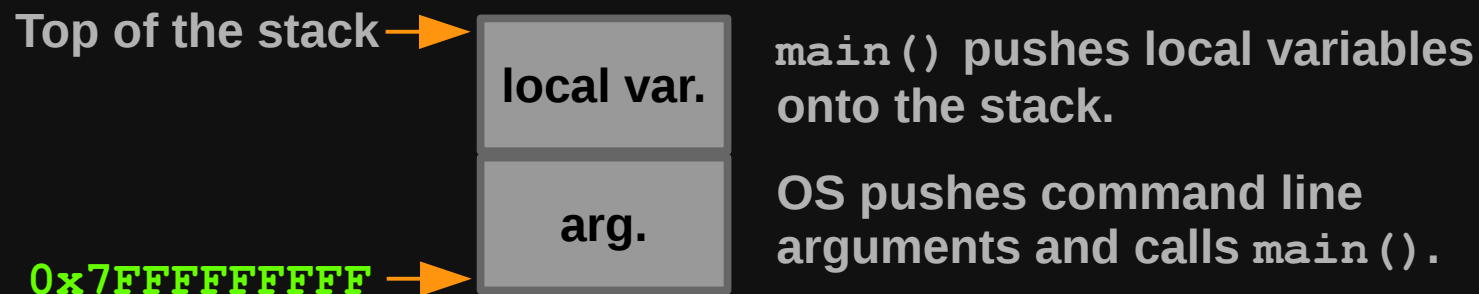
DIY!



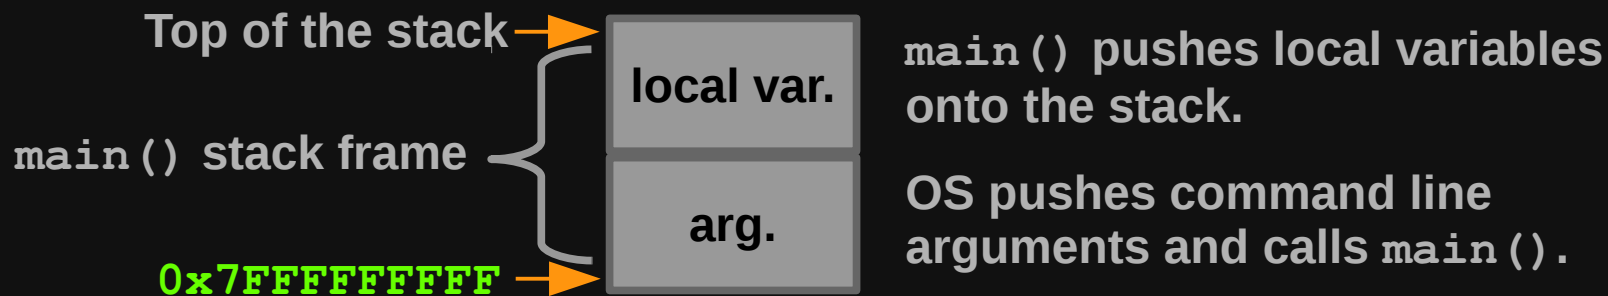
Understanding the Stack frames



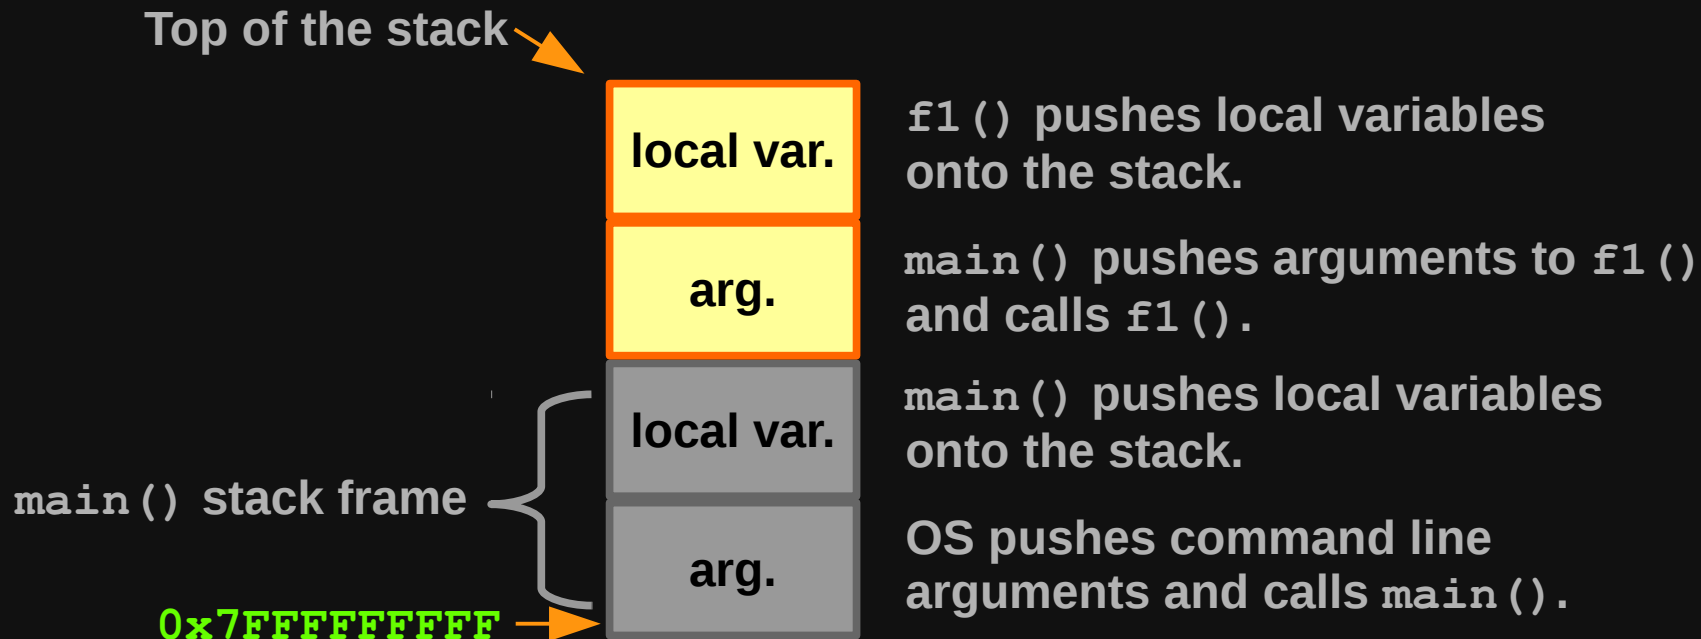
Understanding the Stack frames



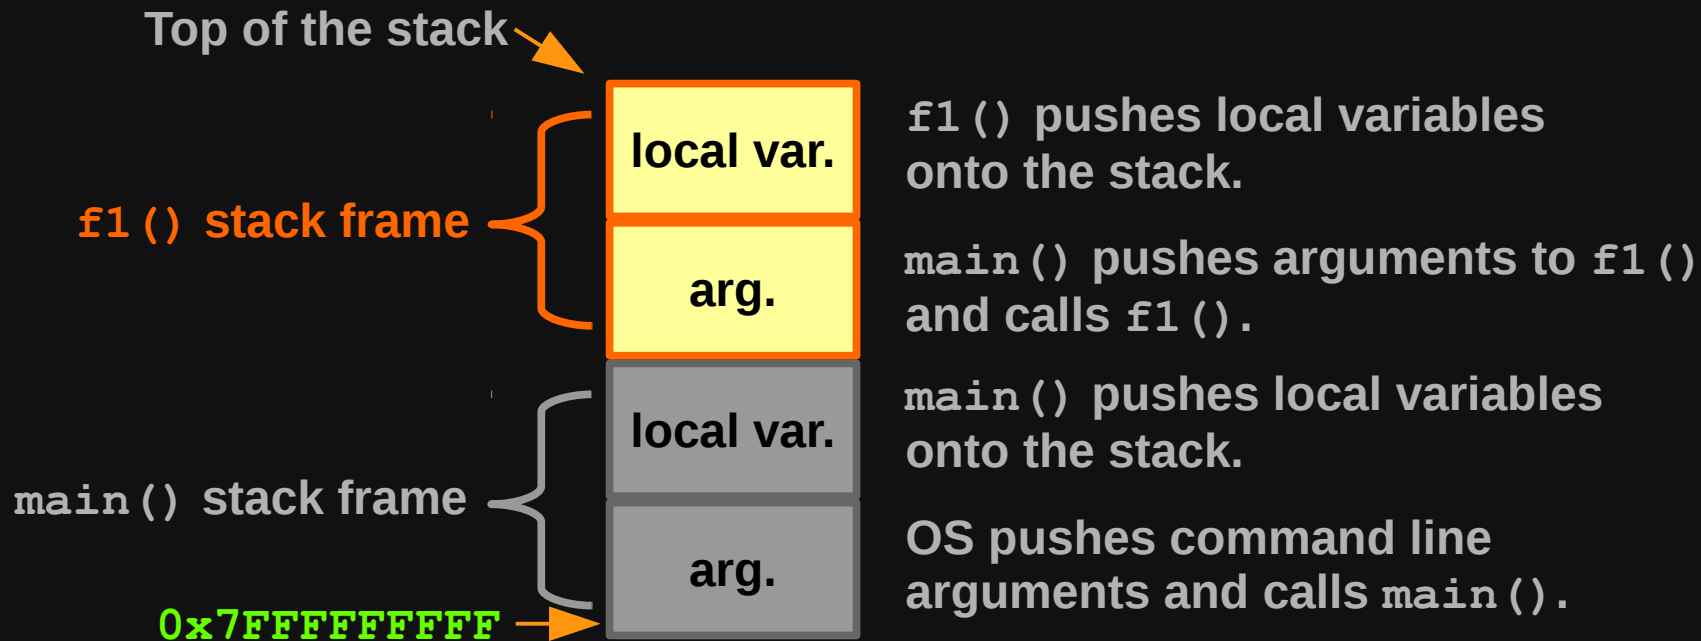
Understanding the Stack frames



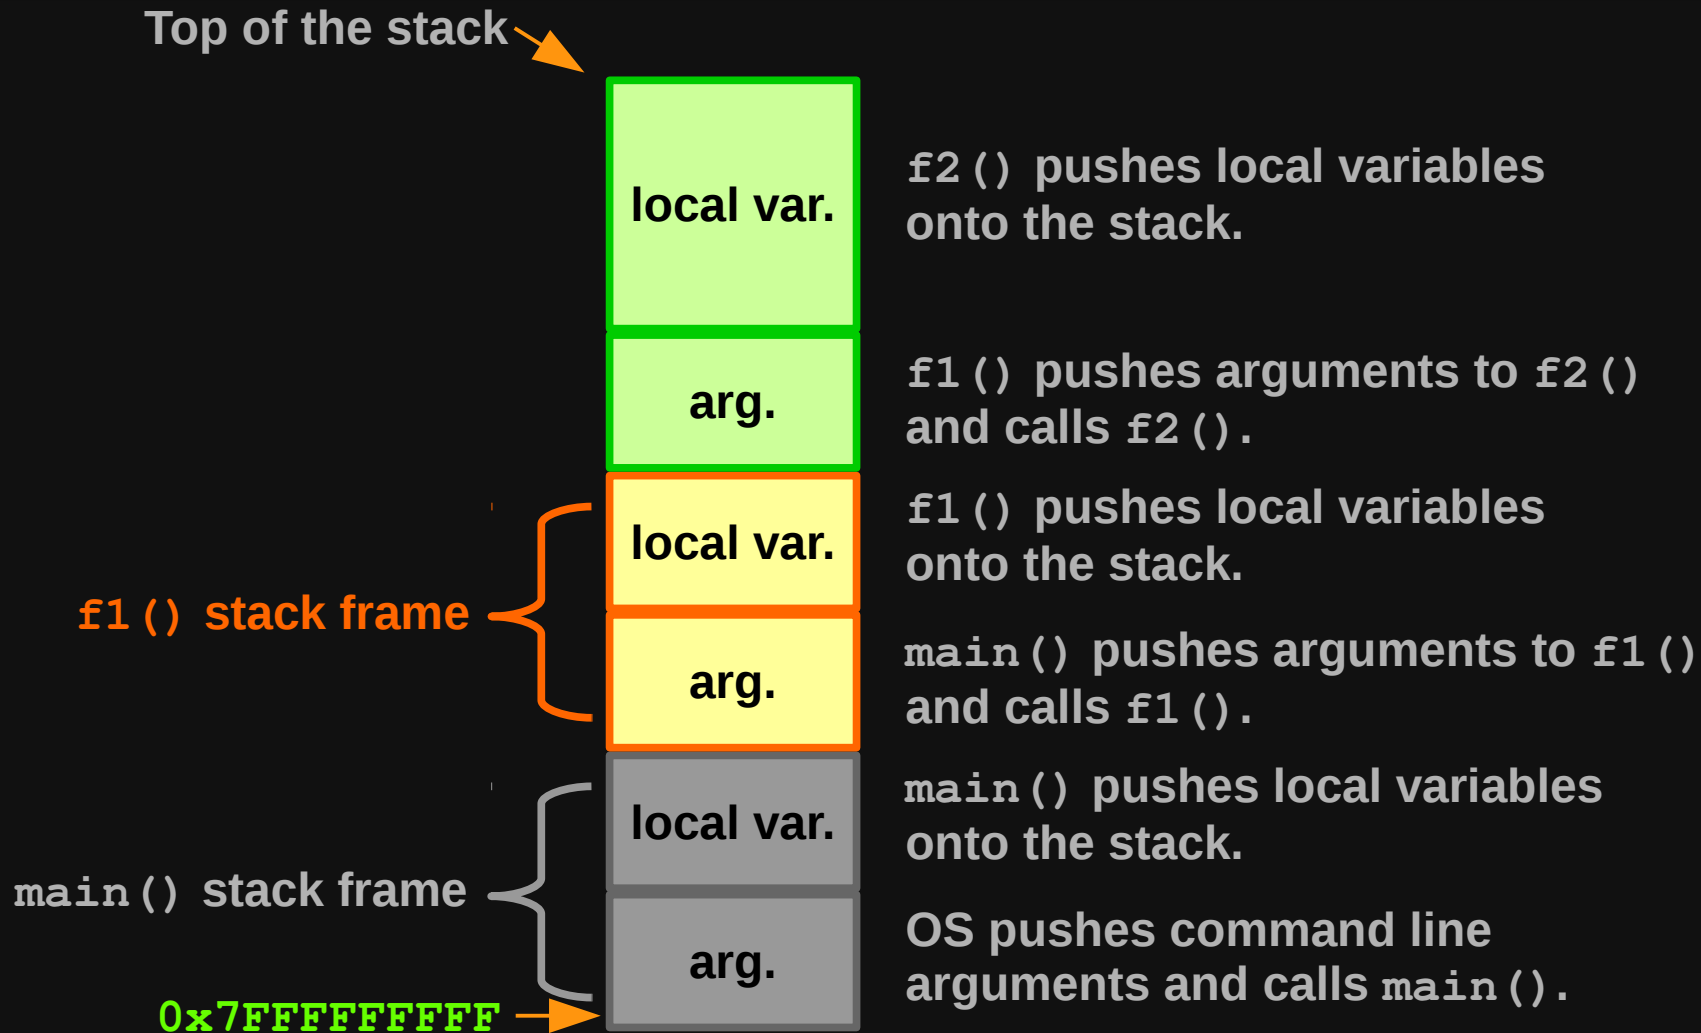
Understanding the Stack frames



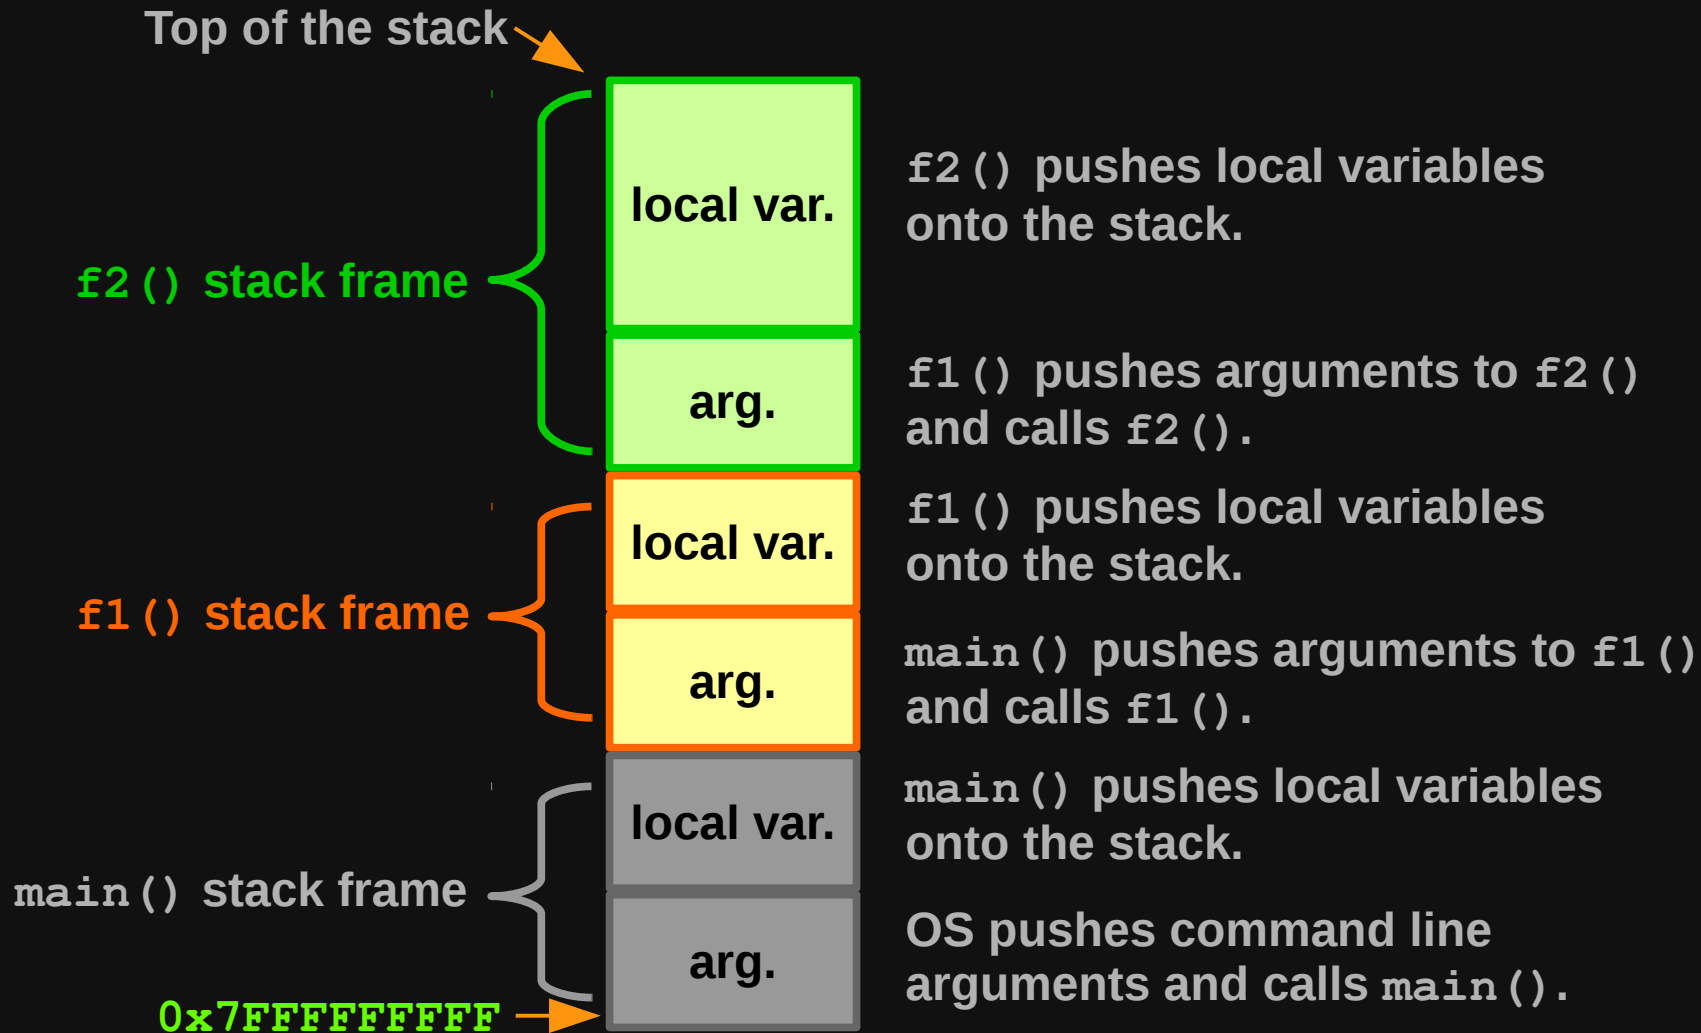
Understanding the Stack frames




Understanding the Stack frames



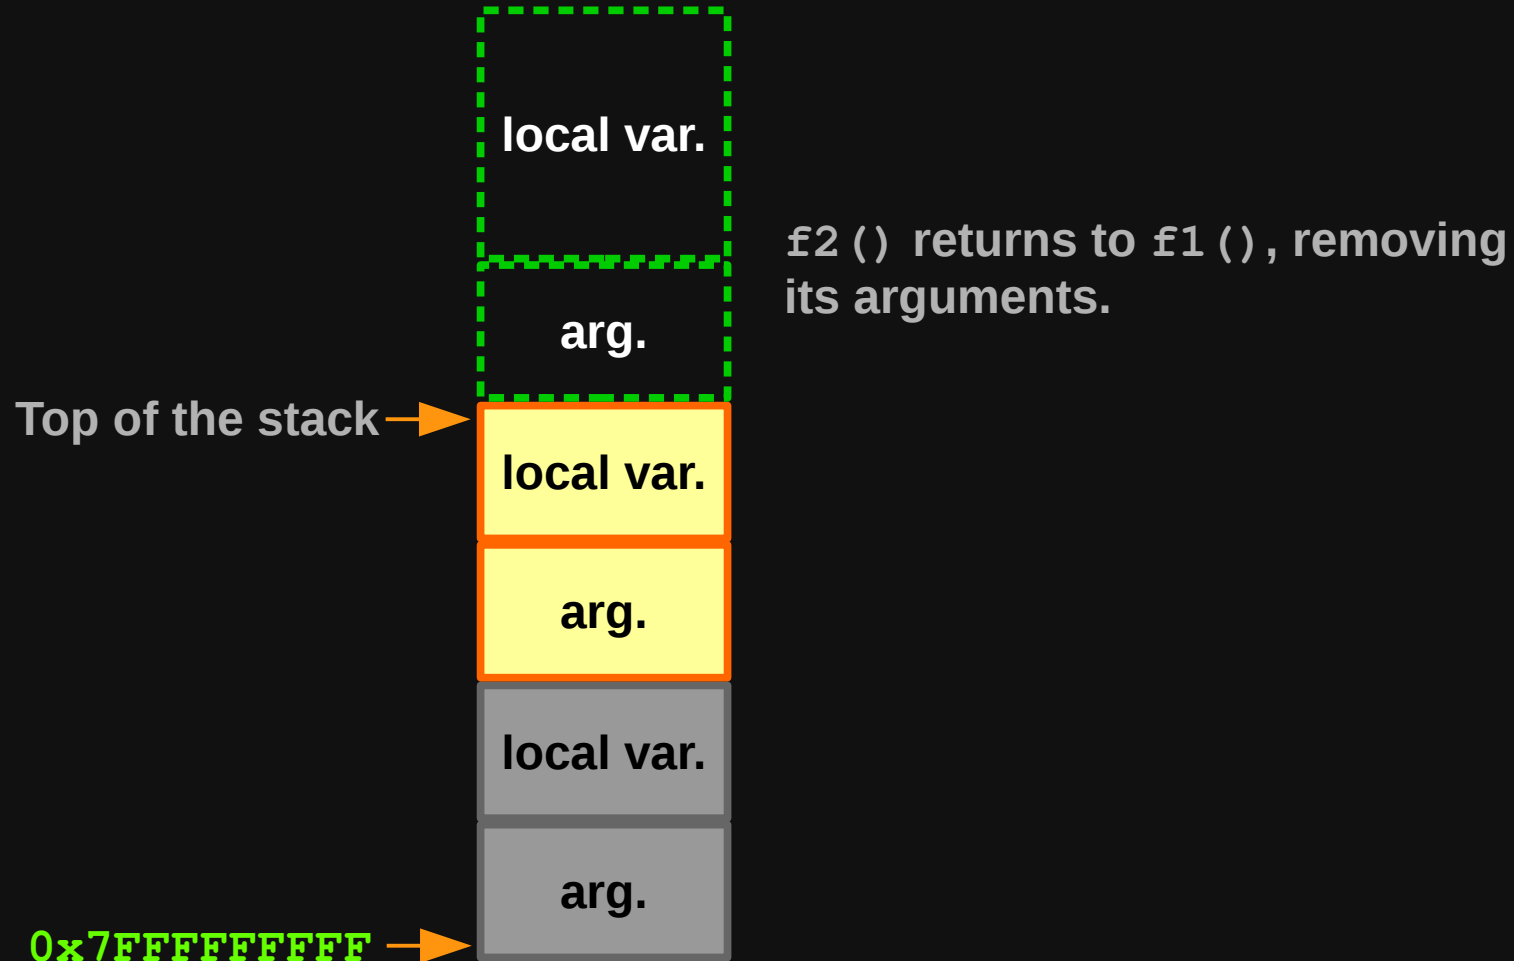
Understanding the Stack frames



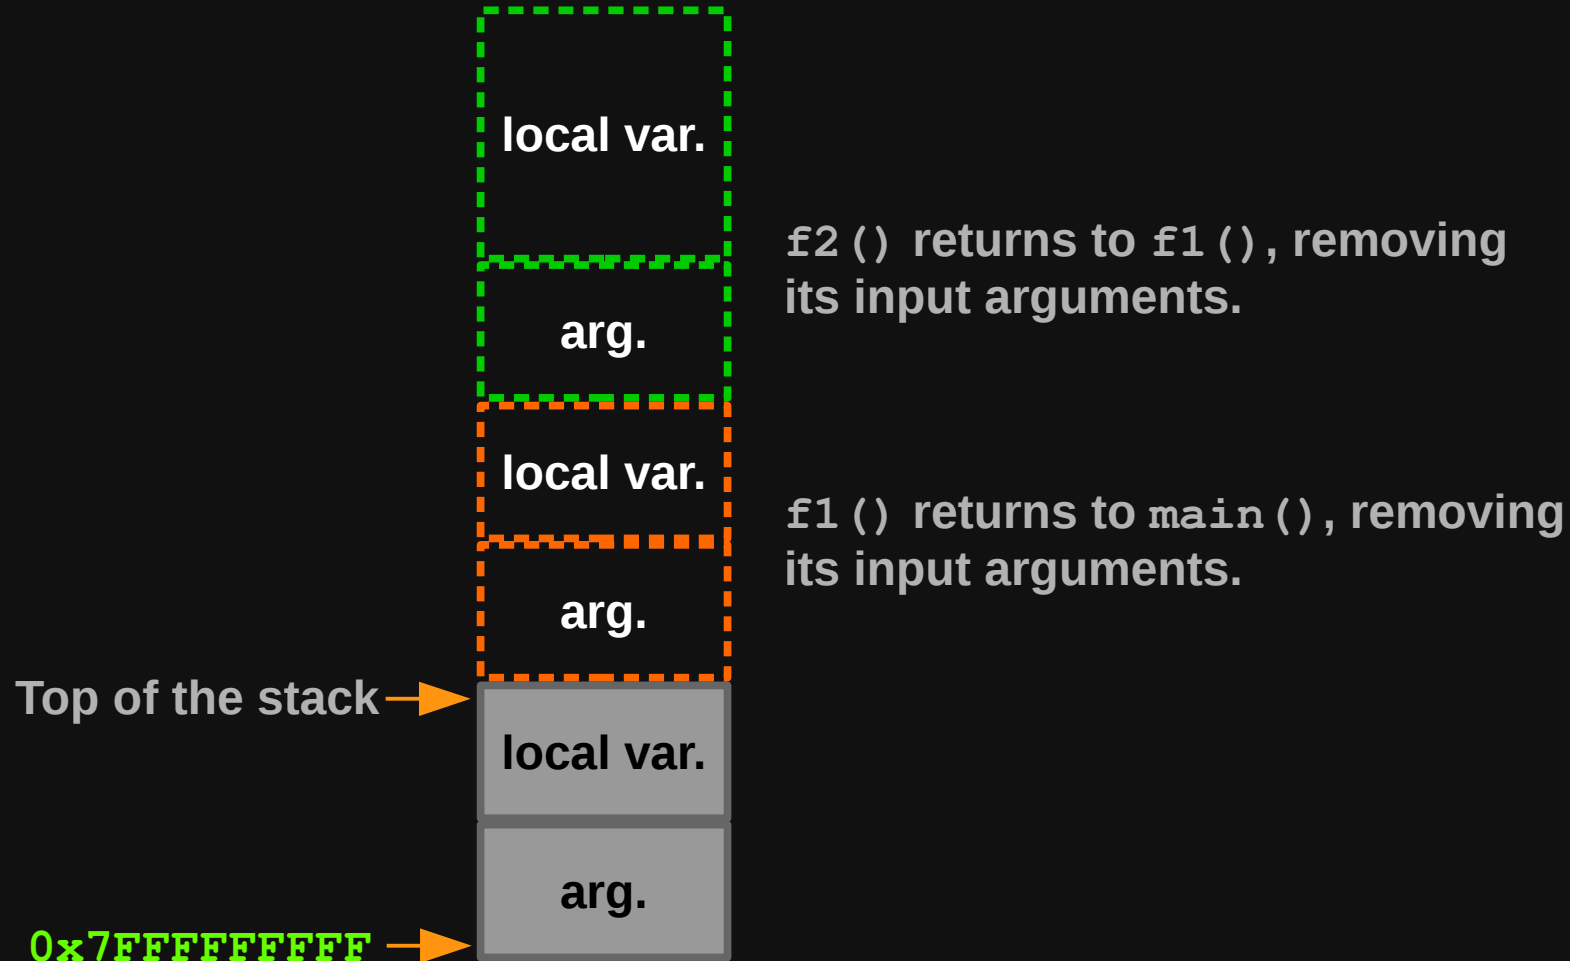
Understanding the Stack frames



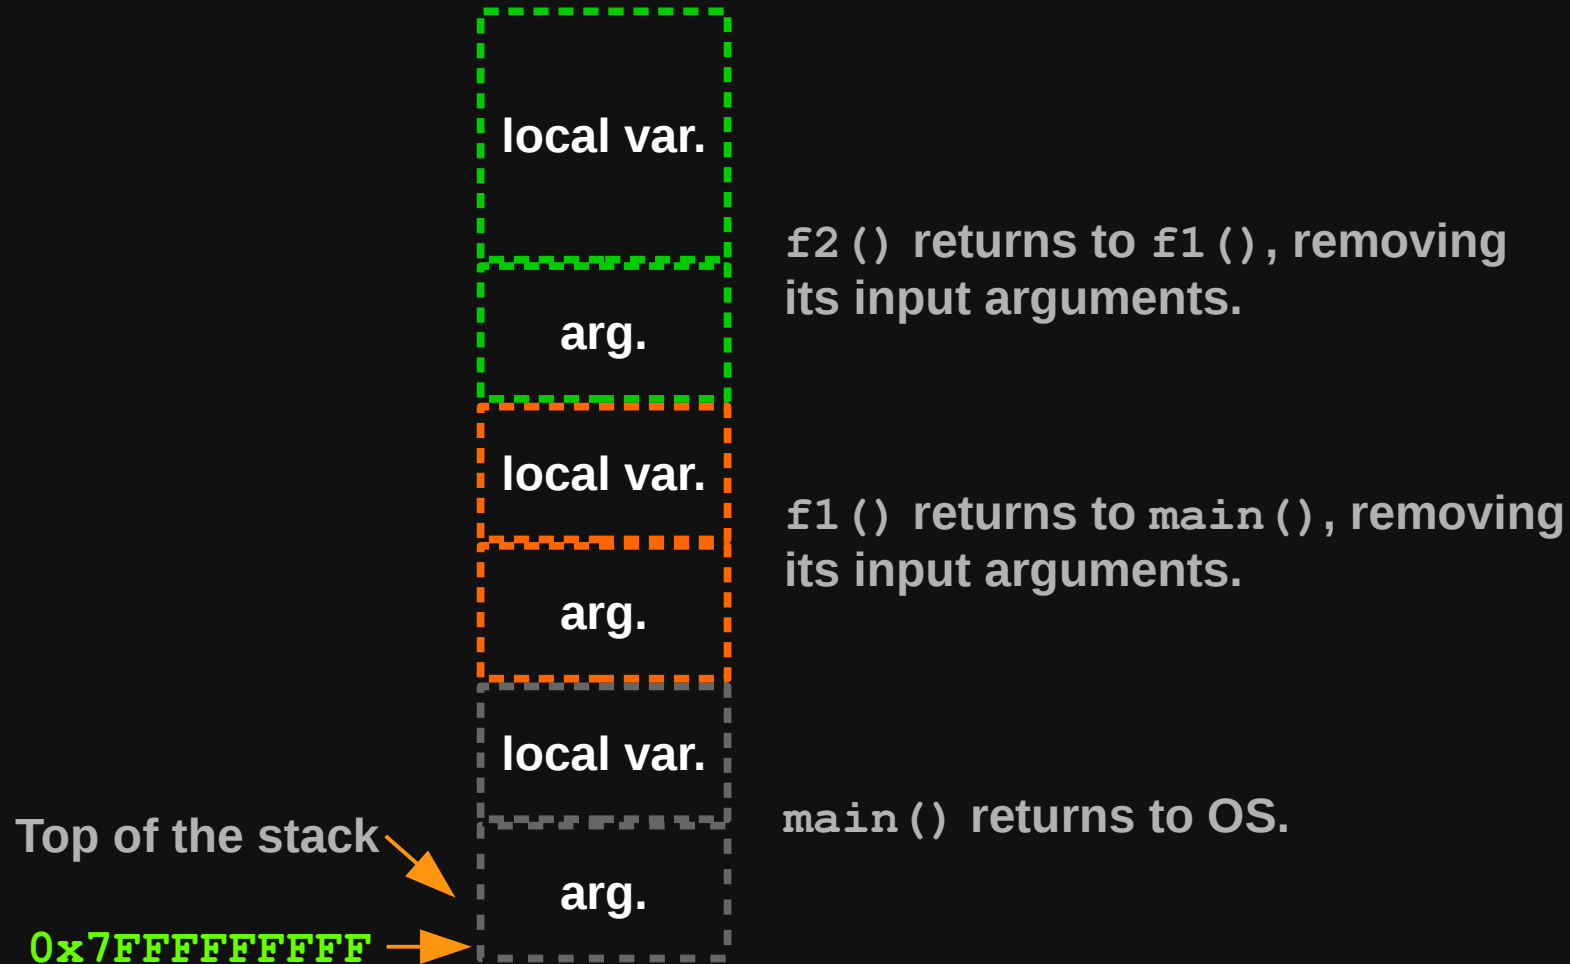
Understanding the Stack frames



Understanding the Stack frames



Understanding the Stack frames

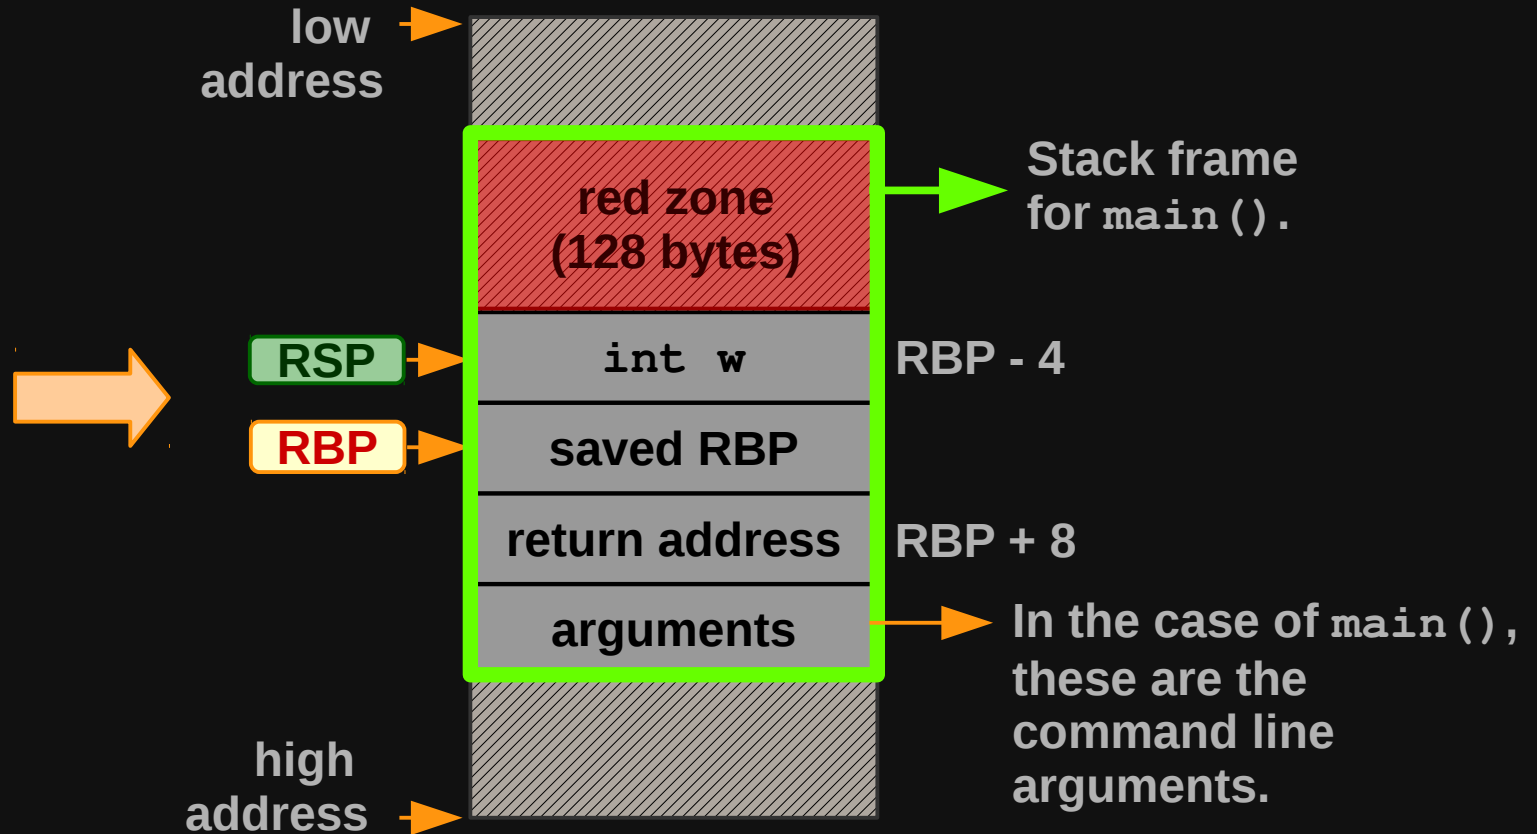


A More Realistic `main()` Stack Frame

`assembly1.c`

```
int x = 10;
int y;

main() {
    int w;
    y = 20;
    w = x + y;
    return 0;
}
```



For those that want to understand all the nitty-gritty about **stacks**, **parameter passing** and **etc.** on the **x86-64**, check the **System V ABI** spec!



Application Binary Interface

- An ABI specifies:
 - Calling convention.
 - Object file formats.
 - Executable file formats.
 - Dynamic linking semantics.
 - Etc.
- Linux uses the **System V ABI**
 - The Executable and Linkable Format (ELF), which is used by Linux, is part of the System V ABI.



System V ABI for x86-64

- The stack grows towards **lower addresses**.
- Parameters to functions are passed in the **registers**:
 - **rdi, rsi, rdx, rcx, r8, r9** (further values are passed on the stack in reverse order).
- **call** instruction pushes the **return address** to the **stack**.
- **ret** instruction pops the **return address** from the **stack** and **jump** to it.
- **Return value** is stored in **rax** register.
- Etc.



Inspecting Stack Frames with GDB

- GDB has some commands for the **inspection** of the **current** stack frame:
 - `info frame`
 - `info args`
 - `info locals`
- To **change** the stack frame:
 - `frame <number>`
- To print the **current** value of the **stack pointer**:
 - `p/x $esp`



Exercise: Inspecting the Stack

assembly2.c

```
int x = 10;

int f2( int c ) {
    int d = c + 100;
    return d;
}

int f1( int a ) {
    int b = a + f2( a );
    return b;
}

int main() {
    int w = f1( x );
    return 0;
}
```

DIY!

Inspect the stack with the following commands:

- **info locals**
- **info args**
- **info frame**

Draw a **diagram** that depicts the **distribution** of data **within** the stack

Experiment the above procedures with **different function parameters!**

- More than 6 parameters.
- Parameters by value and reference
- Structs
- Arrays....

